

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.datasets import make_classification

dataset = pd.read_csv('sonar.csv')

X = dataset.iloc[:, :-1]
y = dataset.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

```

Split the dataset into Training and Test groups (use 20-80 split, i.e. 20% of data will be used for the Test group and 80% for training).

```

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)

```

- Logistic Regression
- Decision Tree Classifier
- Random Forest Classifier (with nb\_trees = 10)
- K- Nearest Neighbors (K-NN)
- Support Vector Machine (SVM) – use the 'rbf' kernel

Train each of the above models and make predictions.

```

logistic_regression = LogisticRegression(random_state=42)
decision_tree = DecisionTreeClassifier(random_state=42)
random_forest = RandomForestClassifier(n_estimators=10,
random_state=42)
knn = KNeighborsClassifier()
naive_bayes = GaussianNB()
svm = SVC(kernel='rbf', random_state=42)

logistic_regression.fit(X_train_scaled, y_train)
decision_tree.fit(X_train_scaled, y_train)

```

```

random_forest.fit(X_train_scaled, y_train)
knn.fit(X_train_scaled, y_train)
naive_bayes.fit(X_train_scaled, y_train)
svm.fit(X_train_scaled, y_train)

y_pred_logistic_regression =
logistic_regression.predict(X_test_scaled)
y_pred_decision_tree = decision_tree.predict(X_test_scaled)
y_pred_random_forest = random_forest.predict(X_test_scaled)
y_pred_knn = knn.predict(X_test_scaled)
y_pred_naive_bayes = naive_bayes.predict(X_test_scaled)
y_pred_svm = svm.predict(X_test_scaled)

results = pd.DataFrame({
    'Actual': y_test,
    'Logistic Regression': y_pred_logistic_regression,
    'Decision Tree': y_pred_decision_tree,
    'Random Forest': y_pred_random_forest,
    'K-NN': y_pred_knn,
    'Naive Bayes': y_pred_naive_bayes,
    'SVM': y_pred_svm
})

print(results)

```

	Actual	Logistic Regression	Decision Tree	Random Forest	K-NN	Naive Bayes
161	Mine	Mine	Mine	Mine	Mine	Mine
15	Rock	Rock	Mine	Mine	Rock	Rock
73	Rock	Mine	Rock	Rock	Rock	Rock
96	Rock	Rock	Rock	Rock	Rock	Rock
166	Mine	Mine	Mine	Mine	Mine	Mine
9	Rock	Mine	Mine	Rock	Rock	Rock
100	Mine	Mine	Mine	Mine	Mine	Mine
135	Mine	Mine	Mine	Mine	Mine	Mine
18	Rock	Rock	Rock	Rock	Rock	Rock
148	Mine	Rock	Mine	Mine	Rock	Rock
171	Mine	Mine	Mine	Mine	Mine	Mine
30	Rock	Rock	Mine	Rock	Rock	Rock

Rock					
155	Mine	Rock	Rock	Rock	Mine
Rock					
180	Mine	Mine	Mine	Mine	Mine
Mine					
125	Mine	Mine	Mine	Mine	Mine
Mine					
197	Mine	Mine	Mine	Mine	Mine
Rock					
164	Mine	Mine	Mine	Mine	Mine
Rock					
190	Mine	Mine	Mine	Mine	Mine
Rock					
84	Rock	Rock	Rock	Mine	Rock
Rock					
75	Rock	Rock	Rock	Rock	Rock
Rock					
124	Mine	Mine	Mine	Mine	Mine
Mine					
170	Mine	Mine	Rock	Rock	Mine
Rock					
104	Mine	Rock	Mine	Mine	Mine
Mine					
101	Mine	Mine	Mine	Mine	Mine
Mine					
69	Rock	Rock	Rock	Rock	Rock
Rock					
25	Rock	Rock	Rock	Rock	Rock
Rock					
95	Rock	Rock	Mine	Rock	Rock
Rock					
16	Rock	Rock	Mine	Mine	Mine
Mine					
141	Mine	Mine	Mine	Mine	Mine
Mine					
185	Mine	Mine	Mine	Mine	Mine
Mine					
154	Mine	Rock	Rock	Mine	Rock
Rock					
68	Rock	Rock	Rock	Rock	Rock
Rock					
66	Rock	Rock	Rock	Rock	Rock
Rock					
120	Mine	Mine	Rock	Mine	Mine
Rock					
147	Mine	Mine	Rock	Mine	Mine
Mine					
98	Mine	Mine	Mine	Mine	Mine
Mine					

138	Mine	Mine	Mine	Mine	Mine
Mine					
167	Mine	Mine	Rock	Rock	Rock
Rock					
45	Rock	Rock	Rock	Rock	Rock
Rock					
113	Mine	Rock	Mine	Mine	Mine
Rock					
65	Rock	Rock	Rock	Rock	Rock
Rock					
178	Mine	Rock	Rock	Rock	Mine
Rock					

	SVM
161	Mine
15	Rock
73	Rock
96	Rock
166	Mine
9	Rock
100	Mine
135	Mine
18	Rock
148	Mine
171	Mine
30	Rock
155	Rock
180	Mine
125	Mine
197	Mine
164	Mine
190	Mine
84	Rock
75	Rock
124	Mine
170	Mine
104	Mine
101	Mine
69	Rock
25	Rock
95	Rock
16	Mine
141	Mine
185	Mine
154	Rock
68	Rock
66	Rock
120	Mine
147	Mine

```
98 Mine
138 Mine
167 Rock
45 Rock
113 Mine
65 Rock
178 Rock
```

Create and print the Confusion Matrix and the accuracy scores for each model.

```
def print_evaluation(model_name, y_true, y_pred):
    print(f"==== {model_name} =====")
    cm = confusion_matrix(y_true, y_pred)
    print(f"Confusion Matrix:\n{cm}")

    acc_score = accuracy_score(y_true, y_pred)
    print(f"Accuracy Score: {acc_score:.4f}\n")

print_evaluation("Logistic Regression", y_test,
y_pred_logistic_regression)
print_evaluation("Decision Tree", y_test, y_pred_decision_tree)
print_evaluation("Random Forest", y_test, y_pred_random_forest)
print_evaluation("K-NN", y_test, y_pred_knn)
print_evaluation("Naive Bayes", y_test, y_pred_naive_bayes)
print_evaluation("SVM", y_test, y_pred_svm)

==== Logistic Regression ====
Confusion Matrix:
[[20  6]
 [ 2 14]]
Accuracy Score: 0.8095

==== Decision Tree ====
Confusion Matrix:
[[19  7]
 [ 5 11]]
Accuracy Score: 0.7143

==== Random Forest ====
Confusion Matrix:
[[22  4]
 [ 3 13]]
Accuracy Score: 0.8333

==== K-NN ====
Confusion Matrix:
[[23  3]
 [ 1 15]]
Accuracy Score: 0.9048
```

```

===== Naive Bayes =====
Confusion Matrix:
[[16 10]
 [ 1 15]]
Accuracy Score: 0.7381

===== SVM =====
Confusion Matrix:
[[22  4]
 [ 1 15]]
Accuracy Score: 0.8810

```

For each model, use the K-fold cross-validation (use K=10; in python - cv=10).

```

def cross_val_and_print(model, model_name):
    scores = cross_val_score(model, X_train_scaled, y_train, cv=10,
                              scoring='accuracy')

    print(f"===== {model_name} (K-fold Cross-Validation) =====")
    print(f"Accuracy Scores: {scores}")
    print(f"Mean Accuracy: {scores.mean():.4f}")
    print(f"Standard Deviation: {scores.std():.4f}\n")

cross_val_and_print(logistic_regression, "Logistic Regression")
cross_val_and_print(decision_tree, "Decision Tree")
cross_val_and_print(random_forest, "Random Forest")
cross_val_and_print(knn, "K-NN")
cross_val_and_print(naive_bayes, "Naive Bayes")
cross_val_and_print(svm, "SVM")

===== Logistic Regression (K-fold Cross-Validation) =====
Accuracy Scores: [0.76470588 0.52941176 1.          0.82352941
0.76470588 0.70588235
0.8125      0.5625      0.875      0.8125      ]
Mean Accuracy: 0.7651
Standard Deviation: 0.1324

===== Decision Tree (K-fold Cross-Validation) =====
Accuracy Scores: [0.88235294 0.70588235 0.64705882 0.70588235
0.76470588 0.70588235
0.875      0.6875      0.75      0.8125      ]
Mean Accuracy: 0.7537
Standard Deviation: 0.0758

===== Random Forest (K-fold Cross-Validation) =====
Accuracy Scores: [0.82352941 0.76470588 0.70588235 0.82352941
0.88235294 0.76470588
0.6875      0.8125      0.75      0.8125      ]
Mean Accuracy: 0.7827

```

Standard Deviation: 0.0563

===== K-NN (K-fold Cross-Validation) =====

Accuracy Scores: [0.94117647 0.58823529 0.70588235 0.82352941  
0.64705882 0.76470588

0.8125 0.8125 0.875 0.75 ]

Mean Accuracy: 0.7721

Standard Deviation: 0.0997

===== Naive Bayes (K-fold Cross-Validation) =====

Accuracy Scores: [0.82352941 0.64705882 0.88235294 0.58823529  
0.58823529 0.76470588

0.625 0.625 0.8125 0.5625 ]

Mean Accuracy: 0.6919

Standard Deviation: 0.1107

===== SVM (K-fold Cross-Validation) =====

Accuracy Scores: [0.94117647 0.76470588 0.94117647 0.82352941  
0.82352941 0.88235294

0.8125 0.8125 0.8125 0.8125 ]

Mean Accuracy: 0.8426

Standard Deviation: 0.0561

For the chosen best performing model select two hyperparameters you would like to tune.

```
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2,  
random_state=42)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

```
gnb = GaussianNB()
```

```
param_grid = {  
    'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6, 1e-5]  
}
```

```
grid_search = GridSearchCV(estimator=gnb, param_grid=param_grid, cv=5,  
scoring='accuracy')
```

```
grid_search.fit(X_train, y_train)
```

```
print("Best Parameters: ", grid_search.best_params_)  
print("Best Accuracy: {:.2f}".format(grid_search.best_score_))
```

```
Best Parameters: {'var_smoothing': 1e-09}
```

```
Best Accuracy: 0.85
```

After tuning the hyperparameter, we can see that the model performed slightly better, however the value of `var_smoothing` is the same as the default one used -  $1e-09$